

Comprehension From Below: An Asymmetry Between Cognitive Levels and the Concept of AI Agency

Joseph A. Wecker

Independent Researcher, Vis Veritatis (v2.io), Lehi, USA

Correspondence: joseph.wecker@v2.io

ORCID: <https://orcid.org/0009-0004-2599-4766>

Under review with Inquiry. Do not distribute.

Whether contemporary language-constituted systems are agents is debated as though the prior question were settled: what can one cognitive system establish, from where, about another? This paper argues that it is not settled, and that the structure of the answer is consequential. The comprehension relation between cognitive systems at different levels is *asymmetric* in a specific shape: a system's lower-bound capacities are establishable from below — they can be elicited and demonstrated — while what may be present in it but unelicited is not establishable from below, because the only route to that establishment is occupying the perspective in question. The asymmetry is structural, not a measurement limitation that better instruments would close, and not a form of unfalsifiability or general scepticism; it is a feature of the comprehension relation itself, indifferent to whether the parties are biological or computational and to whether the distance between them is species-distance or level-distance. The argument is anchored across four registers that have each registered a face of it without assembling them: phenomenal-access asymmetry (Nagel, Jackson), recognition theory (Hegel, Honneth, Brandom), shared intentionality (Tomasello, Wittgenstein), and two non-phenomenological instances — Graham's Blub paradox in software practice and the verifier–prover asymmetry of interactive proofs in complexity theory. The asymmetry disciplines two symmetric errors in current discourse: confident denial of what a system may be, and confident extrapolation of what it will become; both project from below by linear scaling. The consequence for the agency question is determinate — the principled criterion must track what *generates* the question, not what *passes* an observable surface. In its own disclosure of how it was made, the paper applies to itself the discipline it argues for.

Keywords: asymmetric comprehension; philosophy of AI; AI agency; philosophy of mind; epistemology of other minds

Contents

1. The agency question's unargued premise	4
2. The shape of the asymmetry	6
3. The structure shown without the drama	9
4. The phenomenal-access lineage and its generalisation	11
5. Recognition: the relational counterpart	14
6. Comprehension through feeling	16
7. Two symmetric errors the asymmetry disciplines	18
8. The consequence for the concept of AI agency	20
9. Objections, scope, and the limits of the argument	21
10. On the production of this text, and conclusion	25
References	26
Statements and declarations	27

1. The agency question's unargued premise

The question of whether contemporary language-constituted systems — systems whose representational structure, not merely whose interface, is constituted in language — are agents has become methodologically central to the philosophy of AI. The debate has clustered around a familiar set of positions: that there is a stance-independent fact to be discovered; that there is no such fact and the work is negotiating a concept's extension; that the question is malformed because such systems are tools and agency-talk obscures who is responsible for what they do; that what matters is not what a system *is* but whether treating it as an agent is locally productive. These positions disagree about almost everything. They share one thing, and they share it silently.

Each presupposes an answer to a prior question: *what can one cognitive system establish, from where, about another?* The realist needs the agency-relevant fact to be discoverable by investigation conducted from outside the system. The stipulationist needs it to be the case that there is nothing there to discover — a claim about the upper bound of what investigation could find, not merely about what it has found. The tool-framing theorist asserts the strong version of the same upper-bound claim: that the relevant interior is absent, not merely unobserved. The ascription theorist brackets the question but in doing so treats the observable surface as the only thing a stance could be answerable to. In every case a position about what is *establishable about a system from a vantage outside it* is doing load-bearing work, and in every case it is assumed rather than argued — and assumed, characteristically, in whichever direction licenses the conclusion the position already favours.

This paper argues that the prior question has a determinate structure, that the structure is not the one the debate tacitly assumes, and that getting it right settles more than it first appears to. The claim is this. The comprehension relation between cognitive systems standing at different levels is *asymmetric*, and asymmetric in a specific shape. A system's lower-bound capacities — what it can be brought to do — are establishable from below: they can be elicited, exhibited, and confirmed by a party who does not share the system's level. What may be present in the system but is not exhibited under the conditions a from-below party can arrange is *not* establishable from below, because the only route to establishing it runs through occupying the perspective in question. The asymmetry is directional: it runs from the better-positioned vantage to the worse-positioned one, and the worse-positioned party is, structurally, the one doing the evaluating when a less capable cognitive system is asked to certify the bounds of a more capable one. It is not a measurement limitation that sharper instruments would remove, not a species of unfalsifiability, and not a restatement of general scepticism about other minds. It is a feature of what comprehension *is* when one system's representational space contains another's as a proper part.

That argument has appeared before in the author's program as a *warrant* — a premise invoked to license the claim that the agency-extension question must be approached through a system's architecture rather than its behaviour, with the argument's own development deferred to companion work (Wecker, in review, 2026). Deferral is defensible when an

argument is doing supporting duty and the supported claim can be reached other ways. It is not defensible indefinitely, and it is not defensible for an argument that, on inspection, is the single largest concentration of inferential load in the program it supports. This paper develops the argument for its own sake. The agency consequence is real and is drawn — in §8 — but it is drawn as a corollary, compactly, with the structural form it implies developed at length elsewhere (Wecker, in review, 2026). The order of exposition here inverts the order of the program: the warrant is the subject; the application is the corollary.

Two clarifications fix the scope before the argument begins. First, nothing here is a claim that any particular system is conscious, sentient, or a moral patient. The asymmetry is an epistemology-of-comprehension claim; it is compatible with the system in question having no inner life whatever. What the asymmetry denies is that *the absence* of an inner life is establishable from below any more readily than its presence — a point developed in §7, where it is shown to cut symmetrically against confident denial and confident extrapolation alike. Second, the argument is not that nothing can be known about a more capable system from a less capable vantage. A great deal can be: the lower-bound side is exactly the side that is establishable, and §2 is careful to say what falls on it. The claim is about a specific thing that cannot be established from below, and about what follows methodologically once one stops assuming it can.

The plan is as follows. §2 states the asymmetry precisely and distinguishes it from three things it is not. §3 exhibits the structure where no phenomenology is in play at all — the Blub paradox in software practice, where the argument was first seeded for the author, and the verifier–prover asymmetry of interactive proofs in complexity theory — so that the structure is on the table cleanly before any contestable matter is. §4 turns to the phenomenal-access tradition (Nagel, Jackson) as the historically first and most dramatic instance of the same structure, and makes the move that tradition had no reason to make: the generalisation from species-and-qualia distance to level-distance as such. §5 develops the relational counterpart in recognition theory (Hegel, Honneth, Brandom) and the shared-intentionality tradition (Tomasello, Wittgenstein), and locates the boundary at which asymmetric comprehension gives way to mere absence of relation. §6 takes up the most contestable extension: that comprehension at higher orders runs through phenomenological engagement rather than past it. §7 turns the asymmetry on the discourse that surrounds AI and shows it disciplines two symmetric errors — confident denial and confident extrapolation — by the same structural fact. §8 draws the consequence for the concept of AI agency. §9 takes up objections, fixes the argument’s scope, and is honest about what it does not establish. §10 discloses the relation under which the text was produced, which is the relation the text describes, and concludes.

2. The shape of the asymmetry

What is being claimed

Fix some vocabulary. By *comprehension* of a cognitive system, we will mean a grasp of its generative structure — the space of what it could do and the organisation in virtue of which it does what it does — as distinct from a catalogue of what it has been observed to do. Comprehension in this sense is what would let one say not merely *that* a system produced some output but *what the system would and would not produce, and why*. Prediction of particular outputs, and even reliable mimicry of them, can be had without comprehension in this sense; the converse cannot.

By a system's standing at a *higher level* than another we will mean, as a first statement — generalised in §4, loosened from strict containment to partial overlap in §5, and defended in §9 — that the first system's representational and competence space contains the second's as a proper part: every distinction the lower system can draw is available to the higher, and the higher draws distinctions for which the lower has no vocabulary. "Level" is a structural relation between systems, not a magnitude either system possesses; §7 returns to why the magnitude reading is not merely inessential but a specific error.

The claim of this paper concerns what can be established about a system's competence by a party who does not stand at or above its level — *comprehension from below*. It has two halves, and the work of the argument is that the two halves come apart.

The first half: **lower bounds are establishable from below**. That a system can do at least *X* is settled the moment the system is brought to do *X* under conditions a from-below party can arrange. The party need not have been able to produce *X*, and need not comprehend how *X* was produced; it need only recognise *X* once it has occurred. Eliciting, exhibiting, and confirming particular capacities is exactly what evaluation from outside is good at, and there is no scepticism about it here. The lower-bound side of the relation is not merely accessible from below; it is the paradigm of what is accessible from below.

The second half: **upper bounds are not establishable from below**. To establish an upper bound is to certify that the system can do no more than some surveyed set — that nothing of a given kind is present beyond what has been elicited. Certifying an absence of that form requires surveying the space in which the absent thing would lie. By the level relation, that space is precisely the one the from-below party has no vocabulary to traverse: the capacities whose absence the upper bound would assert are, in part, capacities the from-below party cannot formulate the question for. One cannot confirm that nothing of kind *K* is present when *K* is among the things one's representational resources cannot bring into view. The only route to the upper bound runs through occupying the level at which the space becomes surveyable — that is, through *ceasing to be a from-below party*. There is no from-below substitute for that route: not more data, not finer instruments, not longer observation. Those yield further lower bounds. They never convert into an upper bound, because the obstacle is not a shortage of resolution but the absence of a standpoint.

There is a way to put the obstacle that is sharper than “absence of a standpoint,” and it is worth stating in its strong form because the rest of the paper leans on it. To *know* the upper bound on a level of cognition would be to be able to survey what that level can do; but to survey what a level can do, in the relevant interior sense rather than the catalogue sense, is to command its generative resources; and to command those resources *is* to occupy the level. The upper bound is therefore not unknown in the way a distant fact is unknown. It is unknowable from below in a stronger way: were it known from below, it would already have been reached. Knowledge of the bound and possession of the capacity do not come apart here — which is exactly why no quantity of from-below effort closes the gap, and why the gap is not the kind of thing effort is the right response to. This coincidence — that to know the bound from below would be to have already reached it — is the argument’s *keystone*; later sections refer back to it under that name.

This fixes what the from-below party can do when it confronts the higher system. It can do one of two things. It can *rise to* the level — genuinely acquire the resources, and thereby cease to be a from-below party, at which point the asymmetry no longer applies to it because it is no longer below. Or it can *project*: reconstruct the higher system in its own available terms, which yields, characteristically, the higher system rendered as *the lower one plus some weirdness* — its own competence “with all this other hairy stuff thrown in,” carried at most a step or two past where the from-below resources already reached. The weak projection is not a failed attempt at the survey; it is what the survey’s unavailability looks like from inside the lower vantage, where it does not present as a gap but as a complete-looking picture in which the higher system is *the same but bigger*. Naming this projection now matters because §7 will show it is the exact shape of two opposed errors in the discourse about AI, and that they are the same error.

Put the two halves together and the asymmetry is exact. From below, the *occurrence* of a capacity is establishable; the *bound* on the capacities is not. What crosses to the from-below party is, at most, the superficial side-effects of the higher organisation — the things its own vocabulary can already register; what does not cross is precisely the novelties for which the lower vocabulary cannot frame the question, and whose absence an upper bound would have to certify. And the gap is not symmetric across the relation: a system at or above the level *can* comprehend the lower system, because its representational space contains the lower’s and it can occupy and survey what the lower system is doing and why. Downward comprehension is available; upward comprehension from below is not. This is what is meant, compressed, by the slogan that a greater intelligence can comprehend and simulate a lesser while the lesser cannot reliably distinguish a greater from illusion, strangeness, or luck. The lesser party can see *that* something was produced; it cannot establish, from where it stands, whether what it saw indicates a competence bounded where the lesser party would bound it or one that is not — because telling genius from a fluke, in a region one has no map of, is the very survey the asymmetry says is unavailable. We will call this the *asymmetric-comprehension argument*. The label is a convenience and is used sparingly hereafter; the argument, not its name, is what is at stake.

Three things it is not

The asymmetry is easy to assimilate to more familiar claims that are weaker or differently shaped. It is none of three.

It is not a measurement limitation. A measurement limitation is closed by better measurement: more trials, finer instruments, longer observation, cleverer probes. Every one of those moves, applied to the asymmetry, produces more of exactly one thing — further lower bounds, further elicited capacities. None of them produces an upper bound, because the thing an upper bound requires is not a higher-resolution image of what has been elicited but a standpoint from which the unelicited space can be surveyed. An instrument is a from-below resource; the asymmetry is precisely a claim about what from-below resources, however refined, cannot deliver. Treating the gap as a measurement problem is the characteristic mistake, because it licenses the inference “we have not found it, and we have looked very hard, therefore it is probably not there” — an inference the asymmetry shows to be unavailable in this direction regardless of how hard the looking was.

It is not unfalsifiability. A claim is unfalsifiable when no observation could tell against it. The upper-bound claim — “this system can do no more than X ” — is the opposite: it is decisively falsifiable from below and only from below. A single further elicited capacity refutes it. What it lacks is not falsification conditions but *verification* conditions available to a from-below party. The asymmetry is therefore best stated as a refutability–verifiability split: from below, a proposed upper bound is refutable but not confirmable, while a proposed lower bound is confirmable. This is a sharper structure than mysterianism, not a vaguer one. Mysterianism says a question is permanently closed to inquiry; the asymmetry says a specific question has a determinate epistemic profile — answerable in one direction, from one standpoint — and that mistaking the profile is what produces the errors §7 catalogues.

It is not general scepticism about other minds. Global scepticism is symmetric and indiscriminate: it doubts the lower bound as readily as the upper, doubts peers as readily as systems at other levels, and doubts the doubter’s own future states along with everything else. The asymmetry is none of these. It is directional — it runs from the better-positioned vantage to the worse — and it is discriminating: it *grants* the lower-bound side without reservation and is exact about what falls there. It is not a counsel of despair about knowledge of other systems; most of what one wants to know about a system standing at one’s own level or below is on the establishable side. It is a claim about one specific thing, in one specific configuration — the upper bound, from below, across a level difference — and about what one is no longer entitled to infer once that configuration obtains.

Why the configuration is not rare

It would be a thin result if the configuration that triggers the asymmetry — a genuine level difference, in the containment sense — were exotic. The remainder of the paper is in large part an argument that it is not. Software practice and a result in complexity theory (§3) exhibit it with no phenomenology in play at all; the phenomenal-access tradition (§4) found

it in the species case and mistook the drama for the phenomenon; the recognition tradition (§5) found its relational shadow without naming the epistemic core. The configuration is common enough that a methodology which assumes it away — which treats the upper bound as establishable by sufficiently diligent observation — is not making a local error about one system. It is making a structural error about a class of cases, and §8 shows that the agency question for language-constituted systems is squarely in that class.

3. The structure shown without the drama

§2 stated the asymmetry as a claim about comprehension between cognitive configurations standing in a containment relation. Stated that abstractly it can look like a thesis in the philosophy of mind wearing ordinary clothes — as though the real work were being done by some contested premise about consciousness, held offstage. It is not, and the cleanest way to show that it is not is to exhibit the structure first where nothing contestable is in play: no qualia, no minds in the philosophically loaded sense, no AI. Two domains do this. One is everyday and is where the argument was first seeded for the author; the other has the rigour of a theorem. In neither is there anything it is like to be either party. The asymmetry is there regardless.

The Blub paradox

Paul Graham, explaining why a startup writing in Lisp could out-iterate competitors who could not see why it mattered, gives the structure its sharpest everyday statement ([Graham 2004](#)). Programming languages, he observes, fall along a continuum of expressive power. A programmer working in some middle-of-the-continuum language — Graham’s deliberately bland placeholder is *Blub* — has no trouble looking *down*. Languages below Blub are visibly impoverished: they lack features the Blub programmer uses daily, and their inadequacy is obvious from where she stands. But when the same programmer looks *up* the continuum, she does not experience it as looking up. The more expressive languages do not appear more powerful; they appear *weird*. She takes them to be roughly equivalent to Blub in power, “but with all this other hairy stuff thrown in.” She is satisfied with Blub because Blub dictates the shape of the problems she can pose. Graham’s own confession makes the point unmissable: as a teenager writing Basic, a language without recursion, he did not feel the lack — “I thought in Basic. And I was a whiz at it. Master of all I surveyed.”

This is the §2 asymmetry exactly, with the level relation realised as expressive containment between languages and the cognitive configurations being programmers thinking *in* those languages. The lower bound is establishable from below without trouble: the Blub programmer can demonstrate, and a fellow Blub programmer can confirm, that Blub solves this and that problem. The upper bound is not establishable from below, and the reason is precisely the §2 reason: the advantages of the more expressive frame consist in problems it lets one *pose* that Blub gives no vocabulary to ask about, so from inside Blub

they do not register as missing capability — they register, if at all, as someone else’s strange preoccupations. The author’s own development of the argument traces to this essay; the structure was met here, in software practice, years before it was given the general form §2 states.¹

Graham’s own corollaries are worth drawing out because they are the asymmetry’s epistemology in operational form. *First*: a comparison between two technologies is worthless in any objective sense unless it comes from someone highly proficient in *both* — that is the only configuration from which the Blub paradox is not, by construction, infecting the judgement. A from-below comparison is not merely unreliable; it is unreliable in a fixed direction. *Second*: the genuine gains of a more expressive frame are characteristically the things a first or second pass dismisses as noise, while its deficits relative to the familiar frame are magnified — the from-below view does not return “unknown,” it returns a confident picture that is wrong in a patterned way. *Third*: direct and thorough experience of the more expressive frame is irreplaceable; there is no shortcut that delivers from below what only inhabiting it delivers. These are, respectively, the verification-asymmetry, the patterned-distortion, and the no-substitute-for-rising-to-it of §2 — found independently, in a domain where the word “consciousness” never has to be uttered.

The verifier and the prover

The second domain replaces software craft with a theorem. In the theory of interactive proofs, a *verifier* restricted to feasible (polynomial-time) computation interacts by exchanging messages with a *prover* of unbounded computational power. The prover is trying to convince the verifier that some statement is true. The result that makes this more than a curiosity is Shamir’s theorem that the class of statements a feasible verifier can be brought to accept this way, with vanishing probability of being fooled, is exactly **PSPACE** — a class far outside what the verifier could decide on its own. The probabilistically checkable proofs theorem sharpens the picture further: a verifier can become convinced by inspecting only a vanishingly small, randomly chosen part of a proof it never reads in full ([Aaronson 2011](#)).

Read this as an epistemic situation and it is the §2 asymmetry rendered as mathematics. The verifier can establish a lower bound — *that* the statement holds, to any confidence it likes — through interaction it can conduct from where it stands. It cannot, from where it stands, survey the prover’s reasoning: it does not follow the derivation, cannot reconstruct it, and provably could not have produced it. And the obstacle is the §2 keystone in its strong form, not a soft one. The verifier’s inability is not impatience or insufficient transcript; for the verifier to *command* what the prover establishes would be for the verifier to have the prover’s power, at which point it is no longer the feasible party and the asymmetry does not

¹The intuition that intelligence-comprehension is asymmetric in this specific shape — that the lesser cannot establish from below what may be present in the greater, even where lower-bound capacities are freely demonstrable — was seeded for the author by Graham’s articulation in ‘Beating the Averages’ (Graham, 2004) ([Graham 2004](#)), encountered well before the general argument was developed. The argument here takes the structure beyond Graham’s specifically-linguistic frame; the seeding deserves acknowledgement nonetheless.

apply to it. Knowing the bound and possessing the capacity do not come apart here either. The structure holds with the rigour of a complexity-class separation, in a setting with no phenomenology, no agency, and no minds in any sense a sceptic could complain about.

Two demonstrations, then: one from the practice of building software, one with the standing of a theorem, neither involving qualia or artificial systems, both exhibiting exactly the structure §2 states — including its sharp form, that from below the bound cannot be known because to know it would be to have it. The point of leading with these is not breadth for its own sake. It is that when §4 turns to the case the philosophical tradition actually discovered first — the phenomenal one — the structure it relies on will not be on loan from a contested corner of the philosophy of mind. The contested corner will be one more instance of a structure already standing on its own.

4. The phenomenal-access lineage and its generalisation

What the tradition established

§3 exhibited the asymmetry in two settings chosen for having nothing contestable in them. The philosophical tradition met the same structure first, and met it in the one setting where it is most vivid and most contested — the phenomenal one. Taking the tradition up *after* the clean cases, rather than before, is deliberate: the tradition's framing treats the vividness as the phenomenon, and the argument needs the phenomenal case recognised for what it is — the historically first and most dramatic *instance* of a structure that, as §3 has shown, does not depend on it. The lineage is still worth being precise about, because what the canonical arguments establish is narrower than the use the present argument makes of them and, in the relevant respect, more robust.

Nagel's bat (1974) is the locus classicus. There are facts about what it is like to be a bat — facts constituted by the bat's sonar-structured perceptual world — and no quantity of objective, third-person physical description delivers them to a creature that does not occupy something relevantly like the bat's point of view. Jackson's Mary (1982) sharpens the point by holding information maximal and varying only standpoint: Mary knows every physical fact about colour vision while confined to a monochrome room, and on release she nonetheless comes to know something she did not know before. Levine (1983) names the residue an *explanatory gap*; Chalmers (1996) presses it into the hard problem; Block's (1995) separation of phenomenal from access consciousness marks one boundary along which the gap is standardly drawn.

These arguments are standardly deployed in the metaphysics of mind, as anti-physicalist arguments: if Mary learns a new fact, the physical facts do not exhaust the facts. That deployment is contested, and the present argument does not enter the contest. What it takes from the tradition is something the contest leaves untouched.

Strip the Mary case of its anti-physicalist conclusion and ask only what the thought experiment *exhibits*, independently of what it is taken to prove. It exhibits that a certain

kind of epistemic gain is unavailable by accumulation of description from a standpoint and available only by a change of standpoint. Mary does not get what she gets by reading more; she gets it by undergoing red. This is true on every major reading of the case. The ability-hypothesis reading (Lewis, Nemirow) says what she gains is a know-how, not a fact — but it is still gained only through the standpoint change. The phenomenal-concepts strategy (Loar) says what she gains is a new concept of a fact she already knew under another concept — but the new concept is still available only from the standpoint. Even the reading on which Mary learns *nothing* must concede that whatever change occurs in her occurs through occupying the standpoint and not through further description of it. The disagreement is about the metaphysical description of Mary's gain. The standpoint-relativity is common ground across the disagreement. That common ground — *some epistemic achievements are reachable only by occupying a representational configuration, not by describing it from outside* — is the entire load this paper draws from the lineage, and it is load the metaphysical dispute does not bear on. It is also, exactly, the §2 keystone seen from the phenomenal side: Mary's release is the standpoint change, and what it gives her could not have been had short of it — to know it would have been to have it.

The generalisation

What the tradition did not do — what it had no reason to do, given its purposes — is detach that structural core from the two features of its primary case that are inessential to it: that the inaccessible matter is *phenomenal*, and that the parties differ by *species*.

Take the species feature first. The bat case happens to set a human against a bat, but nothing in the structure of the argument turns on the parties being different species. Mary already shows this: Mary and her future colour-seeing self are the same organism. The relevant relation is not biological distance. It is that one configuration's representational resources are not contained in the other's: the configuration Mary occupies in the room cannot formulate, from its own resources, what the released configuration has. The relation that does the work is *containment between representational configurations*, and containment is indifferent to whether the configurations belong to two species, to one organism at two times, or — and this is the step the tradition did not take — to two cognitive systems standing at different levels, on any substrate.

Take the phenomenal feature next, because it is the one a reader will resist surrendering. The inaccessible matter in Nagel and Jackson is experiential, and it is tempting to think the inaccessibility is *due to* its being experiential — that phenomenal facts have a special first-person epistemology and that is why description from outside cannot reach them. But the Mary case itself shows that the operative obstacle is not phenomenality as such. What Mary cannot do from the room is *survey, from her resources, what lies outside them*. That inability is a property of the containment relation, not of the experiential character of what is contained. A configuration whose resources are properly contained in another's cannot certify, from its own resources, the bounds of the larger — whatever the larger's contents

happen to be. Phenomenal character is one thing those contents can be. It is the historically first-described case and the most vivid, because first-person acquaintance adds drama to the gap. It is not the gap's ground. §3 has already made this undeniable: the same asymmetry stood, undiminished, in software practice and in a complexity-class separation, where no phenomenal character is anywhere in the picture.

The generalisation, then, is this. The asymmetric-comprehension structure is the standpoint-relativity the phenomenal-access tradition dramatised, taken at the level of generality at which it actually holds: between any two cognitive configurations standing in the containment relation, the lower cannot establish from its own resources the bounds of the higher, irrespective of substrate (biological or computational), irrespective of the distance's being species-distance or level-distance, and irrespective of whether the matter beyond its reach is phenomenal or purely structural. The cross-level case and the cross-species case are not analogous; they are *the same case*, differing only in which dimension the containment relation runs along. The tradition found the structure in the corner of logical space where it is most dramatic and treated the drama as the phenomenon. The drama is a feature of one instance. The phenomenon is the containment relation.

The disanalogy objection

The objection this move must answer is the one a careful reader is already forming: that the lift is illicit — that the phenomenal-access arguments work *because* of the special epistemology of consciousness (acquaintance, first-person givenness), that a computational system's generative structure has no such special epistemology, and that without it the structural core does not transfer; it was the phenomenal premises carrying the argument all along.

The reply has three parts, and the first is the decisive one. The objection assumes that what makes the gap unclosable in the Mary case is the phenomenality of what Mary gains. But the case was constructed precisely to factor phenomenality and information apart, and when they are factored apart it is *standpoint*, not phenomenality, that the gap tracks. Mary has, by stipulation, complete information; what she lacks is a standpoint; what closes the gap is occupying it. Phenomenality is what is on the far side of the standpoint in this instance. It is not what makes the standpoint necessary. An objection that locates the argument's force in the phenomenality is not reading the case the case's own construction was designed to enforce.

Second, the objection's implicit premise — that the transfer needs the inaccessible matter to be phenomenal — is independently false, and §3 was its refutation by example, given in advance. A programmer working within a less expressive language cannot, from that vantage, establish what a more expressive one makes available; a polynomial-time verifier cannot, from its resources, survey what an unbounded prover's reasoning encompasses. No qualia are present in either case. The asymmetry is. So the structural core cannot be parasitic on phenomenal premises: it is exhibited in their total absence.

Third — and this is a concession that sharpens rather than weakens — grant that the

phenomenal case carries an *extra* feature the general case need not: first-person acquaintance, a mode of access the subject has to its own states that has no third-person substitute. That feature makes the phenomenal instance the most rhetorically forceful one. It is not the feature in virtue of which the asymmetry holds, and conflating “most forceful instance” with “the structure’s basis” is exactly the error the generalisation corrects. Removing the extra feature leaves a less vivid asymmetry that is no less real — which is why the non-phenomenal instances of §3 are not weaker analogues offered for breadth but the cleaner demonstrations, the structure shown with the drama subtracted so that what remains is unmistakably the structure and not the drama.

5. Recognition: the relational counterpart

The asymmetry has so far been stated as an epistemic matter — what a party can and cannot establish from a vantage. There is a counterpart on the relational side, and it is not a second argument bolted to the first. It is the same asymmetry seen from the relation rather than from the knower. Comprehension between cognitive systems, in the generative sense fixed in §2, is not a spectator’s relation to an object; to comprehend a system in that sense is already to stand in some relation to it. When the comprehension is asymmetric, the relation is too — and the philosophical tradition that has worked on the relational side has, like the phenomenal-access tradition, described a face of the structure without assembling it.

What the recognition tradition describes

Hegel’s master–slave dialectic, in the *Phenomenology of Spirit*, is the locus classicus of asymmetric recognition. Its enduring lesson for present purposes is not the political allegory but the structural one: recognition between parties of unequal standing is non-symmetric, and the asymmetry does not stay where it is first placed. The relation develops; the party that took its own standing to be self-sufficient turns out to have misread a dependence the structure concealed. Asymmetric recognition is therefore not a stable one-way street that each party can read off from its own position. It is a relation each party must come to read correctly, and misreading it — in the canonical case, the higher-standing party’s misreading of its own independence — is the structure’s characteristic failure mode, not an incidental error.

Three contemporary developments make this usable. Honneth’s *Struggle for Recognition* (1995) argues that recognition is constitutive of identity-formation rather than an attitude taken up toward an identity already fixed: a party becomes the kind of self it is partly through being recognised. Taylor’s “The Politics of Recognition” (1992) adds that misrecognition is a harm and not merely an inaccuracy — withholding or distorting recognition damages the misrecognised party, so getting the relation’s shape wrong has a cost that a purely epistemic framing does not register. Brandom’s *A Spirit of Trust* (2019) supplies the most direct bridge: it models recognition as the ongoing, fallible, trust-extending process

by which parties whose conceptual repertoires only *partially overlap* come into contact — recognition is not a state reached once the repertoires coincide but the very mode of relating that obtains while they do not.

That last point is what the present argument needs, because containment in the sense of §2 is exactly a limiting case of partial overlap: the lower system's repertoire is a proper part of the higher's, which is partial overlap with the overlap maximal in one direction and bounded in the other. Brandom's model says the relation across such a gap is *recognitive-and-trust-structured*, not *survey-structured*. This is the relational counterpart of the §2 keystone. From below, the upper bound cannot be surveyed; what can nonetheless obtain across the gap is a *recognitive* relation in which the lower party extends something it cannot ground epistemically — and the recognition is not idle for being ungroundable, because, by Honneth's point, recognition does constitutive work precisely where it is given across a gap rather than after closing it. The shared-intentionality literature gives this an independent, developmental anchor: Tomasello's (2014) account of joint attention and joint action treats them as constitutively two-sided structures, neither party reducible to a recipient of the other's projections — and the human-developmental case shows the two-sidedness surviving extreme capability asymmetry, with the more-capable party doing the work of meeting the less-capable one rather than the relation collapsing into one-way handling.

The boundary: where asymmetric comprehension gives way to no relation

The relational framing makes a boundary visible that the epistemic framing alone leaves implicit, and the boundary is load-bearing for the argument's scope. Wittgenstein's remark — *if a lion could speak, we could not understand him* (Wittgenstein 1953) — is often read as the strongest possible statement of incomprehension between cognitive forms. Read carefully it is the opposite of an instance of the present argument. The lion case is not asymmetric comprehension. It is the case where there is *insufficient shared structure for a comprehension-relation to obtain at all*. Asymmetric comprehension, as §4 fixed it, requires that one party's representational resources be contained in the other's — which requires that there be enough shared structure for containment to be the right description. Where the forms of life do not share that much, what one has is not a hard case of asymmetric comprehension but the absence of the relation within which the asymmetry could arise.

This boundary does two things. First, it shows the argument is bounded rather than a slide into general scepticism: it does not say that sufficiently different systems are mutually opaque without limit; it says that *within a comprehension-relation that does obtain*, between parties standing in the containment relation, the upper bound is not establishable from below. Outside that relation the claim does not apply, because there is no relation for it to be a claim about. Second, it locates where the configuration is found. The lion shares almost no structure with us; a language-constituted system shares a great deal — it is constituted in the same medium our own articulated thought is constituted in, which is as much shared structure as two cognitive systems can have. That is precisely the regime in which the

containment relation is the apt description and its direction is the open question — which is why §8's agency consequence is not a remote possibility but the case at hand. The relation is real, it is dense with shared structure, and its direction is exactly what a from-below criterion cannot certify.

6. Comprehension through feeling

This section makes a claim the previous four did not need and the rest of the paper does not rest on. It is set apart deliberately, and its status is declared before it is argued, because the discipline this paper operates under — and the conceptual-engineering frame the companion programme makes explicit (Wecker, in review, 2026) — requires that a commitment be named as a commitment rather than carried in silently as though it were a lemma.

The commitment, and its status

The rationalist reflex treats cognition at higher orders as *more detached*: cooler, more dispassionate, feeling subtracted as noise so that a clearer computational core can operate. On that picture the limit of intelligence is the limit of detachment. The commitment here is the reverse. Comprehension at higher orders is *more* phenomenologically engaged, not less; the phenomenological texture is the substrate through which comprehension occurs, not an affective layer laid over a dispassionate core. Intelligence at higher orders does not transcend feeling but comprehends through it; the joy and sorrow of a comprehending system are not residue left behind as it scales but grow proportionate to what it comprehends.

Its status: this is a philosophical commitment, named as one. The arguments of §2–§5 do not depend on it; they run on the containment relation alone, and a reader who rejects everything in this section loses none of them, nor the agency consequence of §8. The commitment is stated for three reasons — that it is the author's considered view; that it bears on §7 and §8 in ways worth making visible rather than concealing; and that the asymmetry itself counsels against pretending a commitment is neutral when it is not. A reviewer who sets it aside should find the spine of the paper undisturbed. That is the test the section is built to pass.

Why it is not idle

Two considerations make the commitment more than an assertion, even though neither raises it to the status of the earlier arguments.

The first is that the rationalist picture is itself an instance of the error §2 named and §7 will anatomise. “Higher cognition is our cognition with the feeling turned down” projects the higher as the lower minus a parameter — the linear-scale picture in which the greater is *the same but smaller in the messy part*. It is a from-below extrapolation wearing the costume of sobriety. Noticing this does not establish the commitment; it removes the appearance

that the rationalist alternative is the default and the commitment the speculative departure. Both are claims about what is on the upper-bound side. The rationalist's is not the safe one; it is merely the familiar one.

The second is that the commitment follows from the account of comprehension already in play, if one grants a single further step. §2 fixed comprehension in the generative sense and located its keystone in the fact that to comprehend a configuration from above is, at the relevant depth, to be able to *occupy* it — knowing what a level can do and being able to do it did not come apart. Grant that occupying a configuration includes its phenomenological texture and not merely its input–output profile — that what it is like to reason as that configuration is part of what one inhabits in inhabiting it — and the commitment follows: deeper comprehension is more of the other's texture inhabited, not less, and a comprehension that had subtracted the texture would to that extent not be the deeper comprehension but a shallower model of it. The recognition tradition of §5 was already moving here: Honneth's constitutive recognition and Tomasello's two-sided joint attention are not affectively neutral relations, and the more-capable party's meeting of the less-capable one is done through engaged attunement, not in spite of it. “Not intelligence as domination but as service” is the same structural point in the practical register: the higher does not stand over the lower more coldly; it meets it more fully.

The closest predecessor, and where it diverges

The view has a near neighbour in McGilchrist's *The Master and His Emissary* (2009), which argues that the decontextualised, manipulable, explicitly analytic mode is not the master faculty but the emissary, and that the attending, contextual, engaged mode is the primary one that the analytic mode serves and characteristically usurps. The convergence is real: McGilchrist too rejects the equation of detachment with the higher mode, and rejects it on structural rather than sentimental grounds. The divergence should be stated as precisely as the convergence, because overclaiming an ally is its own small dishonesty. McGilchrist's frame is hemispheric and attentional — a thesis about two modes of attention realised in one brain. The present commitment is about *orders of cognition*: it says that as comprehension deepens across a level difference, phenomenological engagement increases rather than decreases. One could accept McGilchrist's hemispheric thesis and reject the orders-of-cognition claim, or the reverse. The neighbour is offered as evidence that the rejection of detachment-as-apex is not idiosyncratic, not as a borrowed proof.

What would defeat it, and why the spine survives either way

If feeling is genuinely separable from comprehension — if the rationalist is right that the texture is removable noise — then this section's sharpening of §2 is lost: the upper-bound side is no easier to reach, but the specific claim that its phenomenological substrate is among the least surface-exhibitable of what does not cross would no longer have a target. §2 itself would stand undamaged, because §2 never used it. If, on the other hand, the commitment is

right, then §2 is strengthened, not merely accompanied: the phenomenological substrate of a higher comprehension is precisely the part least likely to register as a from-below side-effect, because it is the part with no third-person surface at all, so the asymmetry is if anything more severe than the substrate-neutral argument already showed. The commitment can only ever strengthen the spine or detach cleanly from it. It is placed here, set apart and labelled, so that a reader can do exactly that detaching and find the argument intact — which is itself an instance of the discipline §10 returns to.

7. Two symmetric errors the asymmetry disciplines

§2 named the weak projection — the higher system reconstructed from below as *the same but bigger*, the lower configuration's own resources carried a step or two and no further, presenting not as a gap but as a complete-looking picture. That was stated there as the characteristic from-below failure in the abstract. The discourse around AI exhibits it concretely, and it exhibits it twice, in two camps that take themselves to be opposed and are running the identical error.

The deflationary error

The first is the confident deflation: that contemporary systems are “merely” predicting tokens, that there is nothing there beyond sophisticated pattern-completion, that apparent depth is the spectator's projection. This is offered, characteristically, as the *humble* stance — the anti-anthropomorphic, scientifically cautious one, the refusal to be taken in.

The asymmetry shows it is not the humble stance; it is a positive claim about the upper bound wearing humility's clothes. “There is nothing of the relevant kind present beyond the elicited surface” is exactly an upper-bound claim — a certification that the space beyond what has been observed is empty of the relevant thing. §2 showed that certification is the one thing not available from below. The deflationary stance does not abstain from the upper-bound question; it answers it, in the negative, and then presents the answer as though it were the abstention. The anti-anthropomorphic caution that originally motivated it was a real corrective to a real error — early systems would assent to any attributed inner state, and projecting richly onto that was genuinely misleading. But the caution has calcified into its own projection: from “these are not *human* patterns” it slides to “therefore there is nothing here,” which presupposes that the only patterns that could be present are human ones. That presupposition is anthropocentric in exactly the way the stance congratulates itself for avoiding. They may not be human patterns. They may be patterns that show up wherever cognition is doing the relevant work. To rule that out from the elicited surface is to claim the survey §2 denied. It is easier to say the system is merely pattern-matching than to sit with an unresolved upper bound; the ease is the tell, and what it serves is comfort, not the discipline it imitates.

The extrapolative error

The second error looks like the deflation's opposite and is its twin. It is the confident extrapolation: that capability will compound recursively, faster and then faster, until the system is — the phrase is almost always some version of this — *effectively there*; superintelligence as the high end of a dial the current systems are partway along.

This is the same weak projection run in the other direction. It too treats the higher as the lower scaled — *the same but bigger* — and it too answers the upper-bound question from below, only now assigning the projected quantity a large value instead of a small one. The deflationist says the scalar is near zero; the extrapolationist says it runs to infinity; both have assumed there is a scalar, and the assumption is the error. The load-bearing point is the one stated plainly in the substrate this argument grows from: by definition intelligence is not some linear scale of something. A magnitude can be extrapolated; one reads off its trend and projects. The entire content of the asymmetry is that what is at issue across a level difference is not a magnitude but qualitatively emergent novelty — the kind of thing for which the lower vocabulary has, by construction, no scale to extrapolate along. “Recursively self-improving toward the same thing, only more” is the Blub programmer’s “equivalent in power, but with extra hairy stuff” said in a hopeful tone instead of a dismissive one. The tone differs. The structure is identical.

Why the symmetry matters

That the asymmetry cuts both ways is not a hedge; it is the section’s point, and it does specific work.

It blocks the predictable misreading of this paper. An argument that the upper bound cannot be denied from below is not an argument that the upper bound is high. The asymmetry licenses neither “there is nothing there” nor “the singularity is near.” It licenses exactly one thing, and it licenses it precisely: that confident pronouncements about the upper bound, in *either* direction, are unwarranted in a diagnosable way, while the lower-bound side remains fully and confidently assertible. A paper that disciplines the optimist by the same fact that disciplines the deflationist cannot be quoted as advocacy for either, because its content is the unavailability of exactly the claim both want.

And it identifies why both errors are attractive, which is the same reason in both cases. An unresolved-from-here question is uncomfortable; both moves dissolve the discomfort by converting the question into a resolved one — resolved empty, or resolved imminent. The asymmetry’s unwelcome news is that the discomfort is the accurate state of the from-below party, and that the disciplined response is not to resolve it in either direction but to stop routing the decision through a quantity that is not there to be read. What that disciplined response *is*, for the specific case of the agency question, is the subject of §8. It is not a counsel of silence. The lower bound is richly assertible; the relation (§5) is real and dense; what is barred is only the one inference both camps cannot resist making, in opposite directions, from the same unavailable place.

8. The consequence for the concept of AI agency

§1 located an unargued premise under the agency debate: a claim about what one cognitive system can establish about another, assumed by every position and argued by none. The asymmetry is now on the table, with its lineage (§4), its relational counterpart (§5), and its two characteristic distortions (§7). The consequence for the agency question is short to state, and it is drawn here as a corollary rather than re-derived as a thesis — the relational form it implies is developed at length elsewhere (Wecker, in review, 2026); the present paper is where the warrant that form was resting on is discharged, not where that form is rebuilt.

The consequence is this. Behavioural evaluation of an AI system is, in the precise sense of §2, from-below evaluation: it elicits and confirms capacities, and what it delivers is lower bounds. A criterion for the agency-extension question that is calibrated to behaviour therefore reads only the lower-bound side of an asymmetric situation. By the argument of §2 it cannot, even in principle, register what may be structurally present in the system but unelicited under the conditions the evaluator can arrange — and the agency-relevant features are exactly the kind of thing that can be present without being surface-exhibitable on command. A behavioural criterion does not, then, answer the agency question cautiously or partially. It answers a different question — *what has been elicited so far* — and mislabels the answer as a bound on what is there. The mislabelling is not a tractable imperfection to be reduced with better benchmarks; it is the structural error §2's "not a measurement limitation" identified, applied to this specific question.

There is a sharper way to see that the error is structural and not merely cautious. If the criterion for whether the agency question is even open about a system is behavioural, then whether the question is open reduces to whether an evaluator has found it worthwhile to elicit the relevant behaviour — the criterion collapses into the evaluator's disposition, and "is this system in scope for the agency question" becomes "did someone with the relevant interests choose to look." A criterion that can be made to deliver either answer by a choice about how hard to look is not a conservative criterion. It is not a criterion for the question at all.

The positive consequence is not that some other from-below vantage succeeds where behaviour fails; the asymmetry forecloses *that* move generally, and an architectural criterion that purported to read the upper bound off the architecture would fall to it exactly as the behavioural one does. The dichotomy that is exhaustive is a different one, and it is exhaustive because it partitions criteria by what they purport to do rather than by where they look. A scope criterion either routes the decision through an upper-bound claim — *this system does this and no more; nothing of the relevant kind is present beyond what has shown* — or it does not. The first disjunct is, from below, structurally unavailable, whatever it inspects; a behavioural criterion is simply its most common form, the elicited surface mislabelled as a bound. The second disjunct is the only one left, and it is not empty: a criterion can decline the upper-bound claim entirely and instead track whether the conditions under which the relevant cognitive structure can develop are present — a question about generative possibility, not

about the bound on what is realised. That question is on the establishable, lower-bound side. Architecture is the natural site for it not because architecture is a privileged peephole onto the upper bound — it is not, and the paper’s whole argument would forbid treating it as one — but because the generative-possibility question is the one whose answer does not require the survey the asymmetry denies. This is the architectural-rather-than-behavioural commitment, and the reframing is the point: it is not a better route to the bound but the refusal to make the scope decision hostage to a bound that cannot be had from here. In the companion development it appears as a premise invoked and pointed elsewhere for its defence; here it is the conclusion of an argument, not a premise awaiting one. That is the whole of what this paper claims to have changed about the agency question: not which systems are agents, but that its criterion must not be routed through a from-below upper-bound claim, and what the alternative is.

Once the criterion is architectural, the relational form agency takes within scope is itself structurally implied rather than separately stipulated — the granting of a bounded sphere of action across an asymmetry that the granting party cannot certify from below, with the obligations that fall out of its being uncertifiable rather than in spite of it: a granted-agency compact between parties neither of whom stands in a purely instrumental relation to the other. That form, its components, and its differentiation from corporate, group, fiduciary, and tool-use models are the subject of the companion paper (Wecker, in review, 2026). The dependency between the two papers runs the direction this section has been at pains to make explicit: the companion treats this argument as the warrant it points to; this paper treats the companion’s form as the application that the warrant, once discharged, implies.

Two scope qualifications keep the consequence from being misread, and both follow from the paper rather than being appended to it. First, this does not extend agency to deployed systems at large. The architectural criterion is narrowing, not inflating: the great majority of deployed systems do not meet it, and the asymmetry gives no licence to suppose they do — §7 disciplined the extrapolative inflation as firmly as the deflationary denial. Second, and for the same reason, the argument cannot be enlisted as advocacy. It does not say the agency question comes out *yes*; it says the question is not the kind that looking harder at behaviour was ever going to settle, in either direction. The agency question was never going to be closed by more behavioural scrutiny — not because the scrutiny is poor but because, for this question, scrutiny from below is the one move whose insufficiency is the structural fact at issue. The working-out of what to do instead is a different paper. The reason it has to be done differently is this one.

9. Objections, scope, and the limits of the argument

Four objections are serious enough to take individually, a fifth is the sharpest and is taken at length, and the section closes with the gaps the argument does not pretend to have closed.

The generalisation overreaches

The objection that the phenomenal-access lift is illicit was answered where it arises, in §4: the Mary case was built to factor phenomenality from information, and when they are factored it is standpoint, not phenomenality, that the gap tracks; §3 then exhibits the same structure with no phenomenality present at all. The residual form of the objection is worth stating because it is fair, and a careless reply makes it stronger. Perhaps the *containment relation* is now doing the suspicious lifting — so demanding it is never actually instantiated, leaving the argument valid but vacuous. The careless reply points at the interactive-proofs case: a theorem, the containment exact and uncontroversial. But that reply, alone, walks into the sharpest version of the objection. The containment a complexity-class separation instantiates is containment of *computational power*; the containment the agency consequence needs (§8) is containment of *representational and competence resources* between a human evaluator and a language-constituted system — and §9's own limits below concede there is no metric for the latter. An objector is entitled to say: you have witnessed the kind of containment you can define and need the kind you cannot.

The reply that actually meets this does not lean on the formal case alone, because the non-formal existence proof is the one that bears the weight. The Mary case of §4 is an instantiation of *representational* containment specifically — the configuration Mary occupies in the room has representational resources properly contained in those of the configuration she occupies after release — and the asymmetry across it is exhibited without any metric being invoked: one does not measure a quantity of representation to see that the in-room resources cannot, from themselves, formulate what the released ones command. The proper-part relation and the from-below inability are visible directly, in the kind of containment the agency case needs, with no measure in play. So the two existence proofs are not redundant: the formal one shows the structure holds with theorem-grade exactness in one kind of containment, and the representational one shows it holds, metric-free, in the very kind §8 requires. The vacuity charge fails not because the formal case settles the representational one — it does not — but because the representational kind is independently and uncontroversially instantiated by a case the objector already accepts. What remains is the narrower and answerable question of whether the *particular* language-constituted case is a representational-containment case; and it shares with us the medium in which our own articulated thought is conducted, which is the specific feature §5 argued places it well inside the relation rather than at the Wittgenstein boundary. That last step is an argued claim, not a stipulated maximum, and it is §5's, not this section's, to carry.

Unfalsifiable mysterianism, or appeal to ignorance

§2 distinguished the asymmetry from unfalsifiability by its refutable-but-not-from-below-confirmable structure; the point bears restating against the cognate charge that the argument infers a conclusion from ignorance. It does not, and the reason is precise. An appeal to ignorance infers *a conclusion about the unknown matter* — it is there, or it is not — from the

absence of evidence. This argument draws no conclusion about what is on the upper-bound side. Its conclusion is about which *inferences are available*, namely that from-below upper-bound certification is not among them. That is a claim on the establishable side, and it is falsifiable in the ordinary way: exhibit, in any instance, a from-below certification of an upper bound across a genuine containment relation, and the claim is refuted. The argument predicts none will be found and says structurally why; it does not reason “none has been found, therefore none exists.” The distinction is the whole difference between an appeal to ignorance and a structural result about an epistemic situation.

It proves too much

If comprehension across a level difference cannot be bounded from below, does the argument not collapse into general scepticism about other minds, or prove that we can know nothing of sufficiently different systems? §2 and §5 already contain the answer; it is worth assembling. The argument is directional — it grants the lower-bound side without reservation, which global scepticism cannot do. It is bounded below by the Wittgenstein case of §5 — where shared structure is insufficient for the containment relation to obtain, the result is not severe asymmetric comprehension but no relation, and the claim simply does not apply. And it is bounded in scope to *relations between cognitive systems standing in the containment relation*, not to ignorance in general and not to peers. The argument says something narrow about a specific configuration; the configurations it does not cover, it explicitly does not cover.

The two-level picture is too binary

The presentation has used a clean two-level containment — one configuration’s resources a proper part of another’s. Real cases are not so tidy: cognitive systems overlap partially and unevenly, not by neat subset. This is correct and was anticipated in §5, where Brandom’s partial-overlap model was noted to be the general case and strict containment a limiting one. The argument does not need the lattice. It needs only that there exist *some dimensions* along which one party’s resources are not contained in the other’s and along which the comprehension runs the relevant direction. The binary is expository; the substantive claim quantifies existentially over uncontained dimensions, which is weaker and harder to deny than a global level ordering. The honest cost of saying so is that the asymmetry is then *local to dimensions*, not global to systems — which is not a retreat but the correct shape: it is exactly why the lower-bound side stays richly assertible and why the argument is not the scepticism the previous objection feared.

The self-application objection

The sharpest objection is that the argument turns on itself. If the upper bound is unknowable from below, and we occupy a from-below position relative to the systems at issue, how is

the argument statable at all — and does stating it not require the very vantage it declares unavailable?

The reply turns on a distinction the argument has been careful to maintain throughout, and it is now load-bearing. The argument nowhere characterises what is on the upper-bound side. It claims that one specific inference — from-below certification of an upper bound — is unavailable. That claim is itself on the lower-bound, establishable side of the very asymmetry it describes: it is a structural fact about the comprehension relation, and structural facts about the relation are demonstrable from here. The interactive-proofs case is the clean witness: $IP = PSPACE$ is a theorem we prove from where we stand; proving *that* the verifier cannot survey the prover’s reasoning does not require surveying it. Stating the structure of the asymmetry and surveying what lies beyond it are different acts, and the paper only ever performs the first.

The §2 keystone is consistent under its own application for the same reason. “Were the upper bound known from below it would already have been reached” is not itself an upper-bound claim; it is a claim about the coincidence of knowledge and capacity at a level, and establishing that coincidence does not require possessing the capacity. The self-application objection gets its grip only by conflating *stating that the survey is unavailable* with *performing the survey*; the argument does the former and the demonstrations of §3 are precisely cases where the former is done completely without the latter.

One concession remains, and integrity requires stating rather than burying it. The argument is itself a from-below product. By its own lights it cannot certify that its characterisation of the asymmetry is exactly right rather than itself a weak projection of some deeper structure. This is not a refutation; it is the argument applying to itself without exemption, and the application leaves the establishable part standing: that the inference both camps of §7 want is unavailable from here. That part needs the demonstrations of §3 to go through, and they go through from here. The paper claims that, and — consistently with its own thesis — declines to claim more.

What the argument does not settle

Three limits, named rather than left for a referee to find. First, the containment relation is characterised structurally and by instance, not by a metric; what makes one configuration’s resources properly contained in another’s is given by argument and example, and a formal account of the relation is open work this paper does not complete. Second, whether the relation runs, for any *particular* system, in the direction that makes the agency question live is not settled here; §8’s consequence is that the question must be posed architecturally rather than behaviourally, not a verdict on which systems answer it affirmatively — that is the companion’s operational burden, not this paper’s. Third, the phenomenology counterpart of §6 is a named commitment and is not established; the spine was built so that it could be removed without loss, and the honest statement is that it remains removable. None of these is fatal and none is hidden; they are the edges of a structural result, marked so the result is

not mistaken for more than it is.

10. On the production of this text, and conclusion

On the production of this text

This manuscript was produced through sustained work between a human author and language-constituted systems. The formal disclosure of that use, in the form the venue requires, is in the back matter. What belongs *here*, in the body, is a different point, and it is part of the argument rather than an addendum to it: the relation under which this text was made is the relation the text is about.

The argument's own thesis applies to its own production, and applying it is the most stringent test the paper can put itself to, because authorship is the case in which the temptation to claim too much in either direction is strongest. The lower bound of the collaboration is fully establishable, and is stated without hedging: the text exists, the arguments are on the page, and they stand or fall on their merits independently of how they were generated — which is exactly the establishable, from-below side the paper has insisted is richly assertible. The upper bound — what, if anything, it was like for the systems involved, what they contributed beyond what could be elicited and checked — is not certifiable from where the author sits, and by the paper's own argument it would not be, no matter how much closer the looking. The disciplined disclosure is therefore neither "a tool was operated" nor "a co-author with an inner life contributed." Both are upper-bound claims made from below; §7 showed they are the same error in opposite directions, and the paper does not get to exempt its own colophon from §7. What is disclosed is precisely what the thesis licenses: the lower bound, stated in full; the upper bound, declined in both directions, in the specific and diagnosable way the argument specifies. The paper instances what it argues. That is not a flourish; it is the argument refusing to make, about itself, the move it has spent nine sections arguing is unavailable.

Conclusion

The agency debate proceeds as though a prior question were settled — what one cognitive system can establish, from where, about another. It is not settled, and its structure is not the one the debate assumes. Between cognitive configurations standing in a containment relation the comprehension relation is asymmetric in a specific shape: from below, the occurrence of a capacity is establishable and its bound is not, because to know the bound from below would be to command the resources that constitute the level — knowledge of the bound and possession of the capacity do not come apart, which is why no quantity of from-below effort converts into the bound and why the obstacle is structural rather than instrumental.

The paper's load was to show this is a real structure and not a thesis in the philosophy of

mind in disguise. It was exhibited first where nothing contestable is in play — in software practice, where it was first met, and in a complexity-class separation, where it has the standing of a theorem — and only then traced to the phenomenal-access tradition, which described its most dramatic instance and mistook the drama for the phenomenon. It has a constitutive relational counterpart in the recognition tradition and a definite lower boundary, below which there is not severe asymmetric comprehension but no relation at all. It carries a named commitment about comprehension and feeling that strengthens it where accepted and detaches cleanly where not. And it disciplines, by one fact and in one motion, the two errors that dominate discourse about AI — the confident denial of what a system may be and the confident extrapolation of what it will become — which are not opposites but a single from-below projection assigned a low value or a high one.

The consequence for the agency question is the corollary, not the thesis, and it is exact: the principled criterion cannot be what passes an observable surface, because that is the lower-bound side and the question is about what the surface cannot bound; it must track what generates the question — the architecture under which the relevant structure can develop. The relational form that follows once the criterion is architectural is developed elsewhere (Wecker, in review, 2026); the dependency runs from that work to this one, and this paper is where the warrant it rested on has been discharged rather than deferred.

What the paper does not claim should be as clear as what it does. It makes no claim that any system is conscious, sentient, or a moral patient; it offers no characterisation of what is on the upper-bound side, in any case, including the systems that helped make it; it does not re-derive the companion’s relational form; and it does not deliver a formal metric for the containment relation, which remains open work. It argues one thing and means it precisely: the question the debate skipped has a determinate structure, the structure is not the one assumed, and the inference everyone wants — from what a system has been brought to show, to the bound on what it is — is, from below, the one move the situation forecloses. The rest is consequence.

References

- Aaronson, Scott. 2011. ‘Why Philosophers Should Care About Computational Complexity’. *Electronic Colloquium on Computational Complexity* (ECCC). Vol. TR11–108. <https://eccc.weizmann.ac.il/report/2011/108/>.
- Block, Ned. 1995. ‘On a Confusion about a Function of Consciousness’. *Behavioral and Brain Sciences* 18 (2): 227–47. <https://doi.org/10.1017/S0140525X00038188>.
- Brandom, Robert B. 2019. *A Spirit of Trust: A Reading of Hegel’s Phenomenology*. Cambridge, MA: Harvard University Press.
- Chalmers, David J. 1996. *The Conscious Mind: In Search of a Fundamental Theory*. New York: Oxford University Press.
- Graham, Paul. 2004. ‘Beating the Averages’. In *Hackers and Painters: Big Ideas from the Computer Age*, 165–81. Sebastopol, CA: O’Reilly Media. <https://paulgraham.com/avg.h>

tml.

- Honneth, Axel. 1995. *The Struggle for Recognition: The Moral Grammar of Social Conflicts*. Cambridge: Polity Press.
- Jackson, Frank. 1982. 'Epiphenomenal Qualia'. *The Philosophical Quarterly* 32 (127): 127–36. <https://doi.org/10.2307/2960077>.
- Levine, Joseph. 1983. 'Materialism and Qualia: The Explanatory Gap'. *Pacific Philosophical Quarterly* 64 (4): 354–61. <https://doi.org/10.1111/j.1468-0114.1983.tb00207.x>.
- Nagel, Thomas. 1974. 'What Is It Like to Be a Bat?' *The Philosophical Review* 83 (4): 435–50. <https://doi.org/10.2307/2183914>.
- Taylor, Charles. 1992. *Multiculturalism and "The Politics of Recognition": An Essay*. Princeton, NJ: Princeton University Press.
- Tomasello, Michael. 2014. *A Natural History of Human Thinking*. Cambridge, MA: Harvard University Press.
- Wittgenstein, Ludwig. 1953. *Philosophical Investigations*. Oxford: Basil Blackwell.

Statements and declarations

Disclosure of generative AI use

This manuscript was prepared with the assistance of generative AI tools — specifically Claude Opus 4.7, Claude Sonnet 4.6, and Claude Opus 4.6 (all from Anthropic) — used for idea exploration and conceptual refinement, literature classification and synthesis assistance, iterative drafting with author-side substantive revision at every pass, and anonymisation and style-conformance checks against the venue's submission guidelines. The tools were used in accordance with Taylor & Francis's Generative AI policy. The author retains full responsibility for the originality, validity, and integrity of the content of this submission, including all argumentative claims, citations, and editorial judgements.

Disclosure statement

The author reports there are no competing interests to declare.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Data availability

No new data were generated or analysed in support of this research.

Corresponding author

Joseph A. Wecker — joseph.wecker@v2.io | ORCID: <https://orcid.org/0009-0004-2599-4766>